


Setup


We evaluated the performance of the system using 5 clients. The clients were run both on a local machine and on a t2.medium AWS EC2 instance. The probability of performing a trade after a lookup was varied from 0% to 80% in steps of 20%. For each configuration, we measured the average latency for lookup, trade, and order detail requests. Experiments were repeated twice: once with the front-end caching enabled and once with caching disabled.

Native Latency vs Probability with Cache and without Cache:



In the native environment, lookup latency was lower when caching was enabled, with relatively stable values across different trade probabilities. Without caching, lookup latency was slightly higher. Trade and order detail latencies were not affected by caching, since they do not interact with the cache layer. Instead, those latencies scaled with the probability.

AWS Latency vs Probability with Cache and without Cache:



On AWS, latencies were higher due to network and virtualization overhead. However, the same trends were observed. Lookup latency with cache remained lower compared to the configuration without cache. Trade and order latencies were increasing with higher probability.

Cache Replacement:

We verified **LRU** behavior using a test case included in the codebase. After loading 5 stocks into the cache, 6th stock 'NFLX' was accessed, which led to **eviction** of 1st stock 'AAPL'. This behavior was confirmed via log entries that stated: 'Could not find AAPL in cache calling catalog microservice.'

```
#test lru cache eviction
def test_frontend_cache_eviction():

    stocks = ["AAPL", "AMZN", "GOOGL", "META", "NVDA", "NFLX"]
    for stock in stocks:
        requests.delete(f"http://{FRONTENDHOST}:{FRONTENDPORT}/delete/{stock}")
    #cache size is 5, add the first five stocks to the cache
    for stock in stocks[:5]:
        requests.get(f"http://{FRONTENDHOST}:{FRONTENDPORT}/stocks/{stock}")
        time.sleep(0.1)
    # get 6th stock to evict the first stock
    requests.get(f"http://{FRONTENDHOST}:{FRONTENDPORT}/stocks/{stocks[5]}")
    time.sleep(0.1)
    # this is should cause cache miss now as the first stock got evicted
    requests.get(f"http://{FRONTENDHOST}:{FRONTENDPORT}/stocks/{stocks[0]}")
    time.sleep(0.5)
    with open(frontend_log, "r", encoding="utf-8") as f:
        logs = f.read()
    #causes eviction
    assert f"Could not find {stocks[5]} in cache calling catalog microservice" in logs
    #confirms eviction
    assert f"Could not find {stocks[0]} in cache calling catalog microservice" in logs
```

```
src > logs > frontend.log
134 Connection: <socket.socket fd=6, family=2, type=1, proto=0, laddr=('127.0.0.1', 8091),
135 GET [Thread-19 (process_request_thread)] is running to serve ('127.0.0.1', 57588)
136 Could not find NFLX in cache calling catalog microservice
137
138 name: "NFLX"
139 price: 80
140 quantity: 100
141 code: 200
142
143 127.0.0.1 -- [06/May/2025 22:38:37] "GET /stocks/NFLX HTTP/1.1" 200 -
144 Handling get stock lookup requests
145 Client address: ('127.0.0.1', 57591)
146 Connection: <socket.socket fd=6, family=2, type=1, proto=0, laddr=('127.0.0.1', 8091),
147 GET [Thread-20 (process_request_thread)] is running to serve ('127.0.0.1', 57591)
148 Could not find AAPL in cache calling catalog microservice
149
150 name: "AAPL"
151 price: 55
152 quantity: 80
153 code: 200
```

Testing:

```
logs > order_2.log
1 Replicas of (Order 2): are [(1, 'localhost', '8093'), (3, 'localhost', '8095')]
2 {'TransactionNumber': '577', 'Name': 'BoarCo', 'Type': 'sell', 'VolumeTraded': '1'}
3 (Order 2): This is the starting transaction num: 577
4 (Order 2): Started
5 (Order 2): Syncup Request received from 1
6 received last transaction num 577
7 (Order 2): Notified replica with id 2 of leader with id 3
8 Starting syncup
```

```
logs > order_2.log
582 (Order 2): This is the starting transaction num: 1109
583 (Order 2): Started
584 (Order 2): Received request to replicate order from the leader with id 3
585 (Order 2): Received request to replicate order from the leader with id 3
586 (Order 2): Received request to replicate order from the leader with id 3
587 (Order 2): Received request to replicate order from the leader with id 3
588 (Order 2): Received request to replicate order from the leader with id 3
589 [Thread-1 (write_to_disk)] order written to CSV.
590 (Order 2): Received request to replicate order from the leader with id 3
591 (Order 2): Received request to replicate order from the leader with id 3
592 (Order 2): Received request to replicate order from the leader with id 3
593 Starting syncup
594 (Order 2): Syncup Done and completed writing changes to disk
595 (Order 2): Received request to replicate order from the leader with id 3
```

PROBLEMS	DEBUG CONSOLE	OUTPUT	TERMINAL	PORTS		
● ubuntu@ip-172-31-31-15:~/src\$ ls	Dockerfile.catalog	Dockerfile.order	catalog	commands.txt	docker-compose.yml	frontend
native_build.sh	simulate_crashes.sh					
Dockerfile.frontend	README.md	client	deploy_aws.sh	docker_build.sh	logs	
order	test					
● ubuntu@ip-172-31-31-15:~/src\$ simulate_crashes.sh 2						
simulate_crashes.sh: command not found						
● ubuntu@ip-172-31-31-15:~/src\$./simulate_crashes.sh 2						
Crashing Order Replica 2 on port 8094 (PIDs: 110353)						
Restarting Order Replica 2 on port 8094						

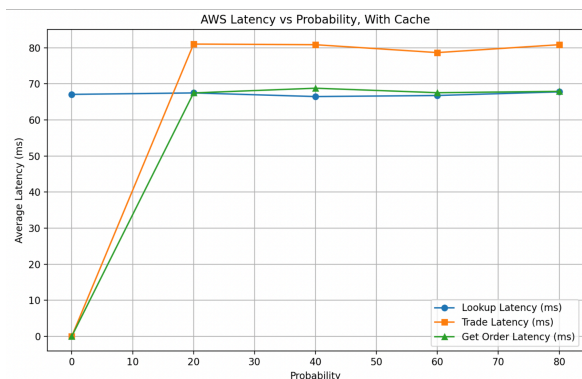
To test fault tolerance, we passed

the replica ID 2 to the script `simulate_crashes.sh`, instructing it to crash Order Replica 2, while clients were actively issuing requests. The system successfully restarted the replica on the same port. After restarting, the replica began synchronizing via the SyncUp protocol. The order logs confirm that it received replicate request from the leader and successfully caught up on missed transactions.

```
logs > order_3.log
6529 (Order 3): Get Order Details Request received
6530 [Thread-1 (write_to_disk)] order written to CSV.
6531 Replicas of (Order 3): are [(1, 'localhost', '8093'), (2, 'localhost', '8094')]
6532 {'TransactionNumber': '1494', 'Name': 'BoarCo', 'Type': 'buy', 'VolumeTraded': '10'}
6533 (Order 3): This is the starting transaction num: 1494
6534 (Order 3): Started
6535 Starting syncup
6536 (Order 3): In Syncup No new entries to write
6537 Replicas of (Order 3): are [(1, 'localhost', '8093'), (2, 'localhost', '8094')]
6538 {'TransactionNumber': '1494', 'Name': 'BoarCo', 'Type': 'buy', 'VolumeTraded': '10'}
6539 (Order 3): This is the starting transaction num: 1494
6540 (Order 3): Started
6541 Replicas of (Order 3): are [(1, 'localhost', '8093'), (2, 'localhost', '8094')]
6542 {'TransactionNumber': '1494', 'Name': 'BoarCo', 'Type': 'buy', 'VolumeTraded': '10'}
6543 (Order 3): This is the starting transaction num: 1494
6544 (Order 3): Started
6545 Starting syncup
6546 (Order 3): In Syncup No new entries to write
6547 (Order 3): Syncup Request received from 2
```

```
PROBLEMS  DEBUG CONSOLE  OUTPUT  TERMINAL  PORTS
```

```
● ubuntu@ip-172-31-31-15:~/src$ ./simulate_crashes.sh 3
Crashing Order Replica 3 on port 8095 (PIDs: 110365)
Restarting Order Replica 3 on port 8095
● ubuntu@ip-172-31-31-15:~/src$ ./simulate_crashes.sh 3
Crashing Order Replica 3 on port 8095 (PIDs: 141581)
Restarting Order Replica 3 on port 8095
● ubuntu@ip-172-31-31-15:~/src$ ./simulate_crashes.sh 3
Crashing Order Replica 3 on port 8095 (PIDs: 141628)
Restarting Order Replica 3 on port 8095
● ubuntu@ip-172-31-31-15:~/src$ ./simulate_crashes.sh 2
Crashing Order Replica 2 on port 8094 (PIDs: 129781)
Restarting Order Replica 2 on port 8094
● ubuntu@ip-172-31-31-15:~/src$ ./simulate_crashes.sh 1
Crashing Order Replica 1 on port 8093 (PIDs: 110338)
Restarting Order Replica 1 on port 8093
● ubuntu@ip-172-31-31-15:~/src$ ./simulate_crashes.sh 3
Crashing Order Replica 3 on port 8095 (PIDs: 141662)
Restarting Order Replica 3 on port 8095
```



order_log_1.csv	order_log_3.csv	order_log_2.csv
order > order_log_1.csv	order > order_log_3.csv	order > order_log_2.csv
1974 1954,BoarCo,buy,3	1974 1955,AAPL,sell,6	1974 1955,AAPL,sell,6
1974 1955,AAPL,sell,6	1975 1956,AMZN,sell,1	1975 1956,AMZN,sell,1
1975 1956,AMZN,sell,1	1976 1957,MenhirCo,buy,1	1976 1957,MenhirCo,buy,1
1976 1957,MenhirCo,buy,1	1977 1958,BoarCo,sell,4	1977 1958,BoarCo,sell,4
1977 1958,BoarCo,sell,4	1978 1959,AAPL,sell,2	1978 1959,AAPL,sell,2
1978 1959,AAPL,sell,2	1979 1960,AMZN,buy,10	1979 1960,AMZN,buy,10
1979 1960,AMZN,buy,10	1980 1961,META,sell,5	1980 1961,META,sell,5
1980 1961,META,sell,5	1981 1962,AAPL,buy,2	1981 1962,AAPL,buy,2
1981 1962,AAPL,buy,2	1982 1963,META,sell,7	1982 1963,META,sell,7
1982 1963,META,sell,7	1983 1964,BoarCo,buy,9	1983 1964,BoarCo,buy,9
1983 1964,BoarCo,buy,9	1984 1965,MenhirCo,sell,2	1984 1965,MenhirCo,sell,2
1984 1965,MenhirCo,sell,2	1985 1966,AMZN,buy,3	1985 1966,AMZN,buy,3
1985 1966,AMZN,buy,3	1986 1967,META,buy,9	1986 1967,META,buy,9
1986 1967,META,buy,9	1987 1968,AAPL,buy,1	1987 1968,AAPL,buy,1
1987 1968,AAPL,buy,1	1988 1969,META,sell,4	1988 1969,META,sell,4
1988 1969,META,sell,4	1989 1970,AMZN,sell,10	1989 1970,AMZN,sell,10
1989 1970,AMZN,sell,10	1990 1971,MenhirCo,buy,7	1990 1971,MenhirCo,buy,7
1990 1971,MenhirCo,buy,7	1991 1972,META,sell,7	1991 1972,META,sell,7
1991 1972,META,sell,7	1992	1992

We tested the failure of the leader “Replica 3”. The front-end re-elected a new leader automatically. Clients continued sending requests without interruption, and no latency spikes were observed in the performance plots. The **final CSV files from all replicas were verified** to be consistent. Together, the script usage, logs, latency graphs, and file comparisons confirm that failure handling was **transparent to clients** and system state remained correct.

AWS Setup

From the AWS academy learner lab, we first download the pem file to access the server using ssh

Command to register public key in the cloud in us-east-1 region

```
chmod 400 labsuser.pem
```

```
ssh-keygen -y -f labsuser.pem > labsuser.pub
```

```
aws ec2 import-key-pair \
```

```
--key-name labsuser \
```

```
--public-key-material fileb:///labsuser.pub
```

```
aws ec2 describe-key-pairs --query "KeyPairs[*].KeyName"
```

Now run the script from the src directory to start EC2 instance (t2.medium) using aws cli commands

```
./deploy_aws.sh
```

Create a tarball from outside src (project root folder) while excluding .venv and other ignored files (pycache) via .gitignore:

```
git archive --format=tar.gz --output=lab3.tar.gz HEAD
```

Transfer Package to AWS EC2

```
scp -i src/labsuser.pem lab3.tar.gz ubuntu@<Public_IP>:~/
```

SSH into EC2 Instance

```
ssh -i stock-key.pem ubuntu@$PUBLIC_IP
```

Extract Project Archive

```
tar -xzf lab3.tar.gz
```

Install Python and Dependencies

- `sudo apt update`
- `sudo apt install python3-pip -y`
- `pip install -r requirements.txt`

Symlink python

```
sudo ln -s /usr/bin/python3 /usr/bin/python
```

Run Server Using Script

```
./native_build.sh
```